

CS 512 Assignment 4: Report

Fall 2018

Herick Jayesh Asmani

A20399752

Framework Used: Keras.

Instructions to run programs:

1. Install various python packages such as numpy, OpenCv, TensorFlow, Keras, Matplotlib.
2. To train the model and visualize various plots (accuracy, loss, precision, recall) run `cnn.py`.
3. To visualize the model and train and evaluation plots, in your command prompt run `tensorboard --logdir ./models`. It will give you a link. Copy and past that link into your web browser to see the visualization.
4. To classify unseen digits images run `cnn_test.py`. and provide path of the image.

Deliverable 1:

Since we are working with MNIST datasets, the output labels are 0 – 9 but we want that given an input image, the classifier should classify the image as either even or odd. So, for the classifier to classify input image as even or odd, we map the output labels to 1s and 0s i.e. map even numbers (0, 2, 4, 6, 8) to 1 and odd numbers (1, 3, 5, 7,9) to 0. Therefore, now we have binary output labels.

Model Details:

Based on the instructions given in the Assignment 4 description, following Model is built.


Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 28, 28, 32)	832
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_2 (Conv2D)	(None, 14, 14, 64)	51264
max_pooling2d_2 (MaxPooling2D)	(None, 7, 7, 64)	0
dropout_1 (Dropout)	(None, 7, 7, 64)	0
flatten_1 (Flatten)	(None, 3136)	0
dense_1 (Dense)	(None, 128)	401536
dense_2 (Dense)	(None, 50)	6450
dense_3 (Dense)	(None, 2)	102
Total params: 460,184		
Trainable params: 460,184		
Non-trainable params: 0		

Result:

After compiling the above model with categorical cross entropy loss (since I am using Softmax output) and gradient descent optimizer with learning rate of 0.001. The model is trained for 5 epochs and following training and test metrics curves are obtained.


Training Accuracy:

acc




Training Loss:

loss




Training Precision:

precision




Training Recall:

recall




Test Accuracy:

val_acc




Test Loss:

val_loss




Test Precision:

val_precision



Test Recall:

val_recall



The values of accuracy, loss, precision and recall for training and test of the above CNN model at the final step is as follows:

```
Epoch 5/5
60000/60000 [=====] - 185s 3ms/step - loss: 0.2078 - acc: 0.9196 - precision: 0.5000 - recall: 1.000
0 - val_loss: 0.1530 - val_acc: 0.9452 - val_precision: 0.5000 - val_recall: 1.0000

Epoch 00005: val_loss improved from 0.19593 to 0.15301, saving model to ./models/model.weights.best.hdf5
Test loss: 0.1530118324995041
Test accuracy: 0.9452
Test precision: 0.5
Test recall: 1.0
```

Deliverable 3:

Algorithm Details:

In order to classify the unseen handwritten digits in an image, we have to do some preprocessing. First, we obtain an image by giving a path of the desired image file. After obtaining the image, we resize the image to match the dimensions of the images on which we trained the CNN model i.e. we resize the image to 28x28. Now the CNN model was trained on binary images, so to classify the unseen image correctly, we first convert the RGB image to grayscale image and then apply Gaussian blur and Adaptive threshold technique to binarize the image. The algorithm calculates the threshold for a small regions of the image. So we get different thresholds for different regions of the same image and it gives us better results for

images with varying illumination. However, it thresholds the image as 0 and 255. Therefore, we binarize it by dividing the result by 255. Now the image is processed i.e. ready for classification. So, now we import the CNN Model we created in cnn.py and load its best weights for classification. Before feeding the image, we do reshaping of the image to (1, 28, 28, 1) because that's the dimension of the input, the CNN accepts and then we feed the image to this model and predict its class as either even (1) or odd (0).

Note: The unseen images used here were obtained by drawing digits on a white background using “Sketches” app from Sony.

Results:

The custom CNN model was able to correctly classify the digits except for some.

9 – correctly classified as odd.




```
Enter q if you want to exit else enter any other key to continue.
Type the Path of your image file: C:\Users\Herick\Desktop\MS\CS512 Computer Vision\Assig
C:\Users\Herick\Desktop\MS\CS512 Computer Vision\Assignments\Numbers\Sketches\Nine.png
True
```

```
Type the Path of your image file: C:\Users\Herick\Desktop\MS\CS512 Computer Vision\Ass:
C:\Users\Herick\Desktop\MS\CS512 Computer Vision\Assignments\Numbers\Sketches\Nine.png
True
```

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 28, 28, 32)	832
max_pooling2d_3 (MaxPooling2)	(None, 14, 14, 32)	0
conv2d_4 (Conv2D)	(None, 14, 14, 64)	51264
max_pooling2d_4 (MaxPooling2)	(None, 7, 7, 64)	0
dropout_2 (Dropout)	(None, 7, 7, 64)	0
flatten_2 (Flatten)	(None, 3136)	0
dense_4 (Dense)	(None, 128)	401536
dense_5 (Dense)	(None, 50)	6450
dense_6 (Dense)	(None, 2)	102
Total params: 460,184		
Trainable params: 460,184		
Non-trainable params: 0		

```
[0]
The digit is odd
```

0 – correctly classified as even.



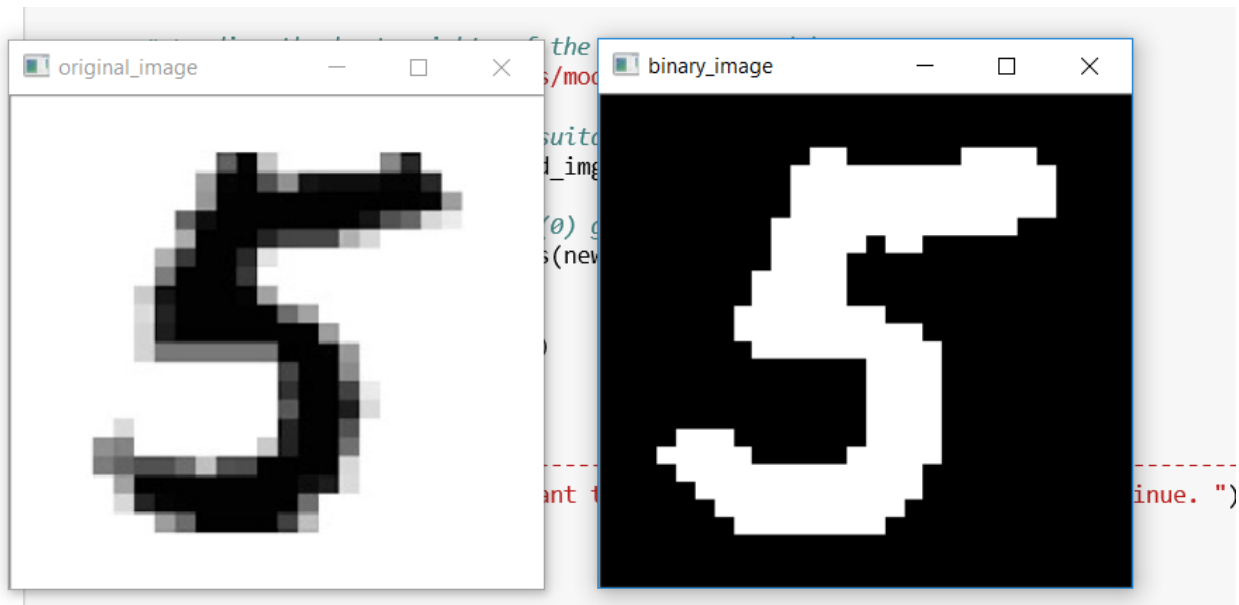
Type the Path of your image file: C:\Users\Herick\Desktop\MS\CS512 Computer Vision\Assignments\Numbers\Sketches\Zero.png
True

Type the Path of your image file: C:\Users\Herick\Desktop\MS\CS512 Computer Vision\Assignments\Numbers\Sketches\Zero.png
True

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 28, 28, 32)	832
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_2 (Conv2D)	(None, 14, 14, 64)	51264
max_pooling2d_2 (MaxPooling2D)	(None, 7, 7, 64)	0
dropout_1 (Dropout)	(None, 7, 7, 64)	0
flatten_1 (Flatten)	(None, 3136)	0
dense_1 (Dense)	(None, 128)	401536
dense_2 (Dense)	(None, 50)	6450
dense_3 (Dense)	(None, 2)	102
=====		
Total params: 460,184		
Trainable params: 460,184		
Non-trainable params: 0		

[1]
The digit is even

5 – correctly classified as odd.



```
] : if __name__ == '__main__':  
    main()
```

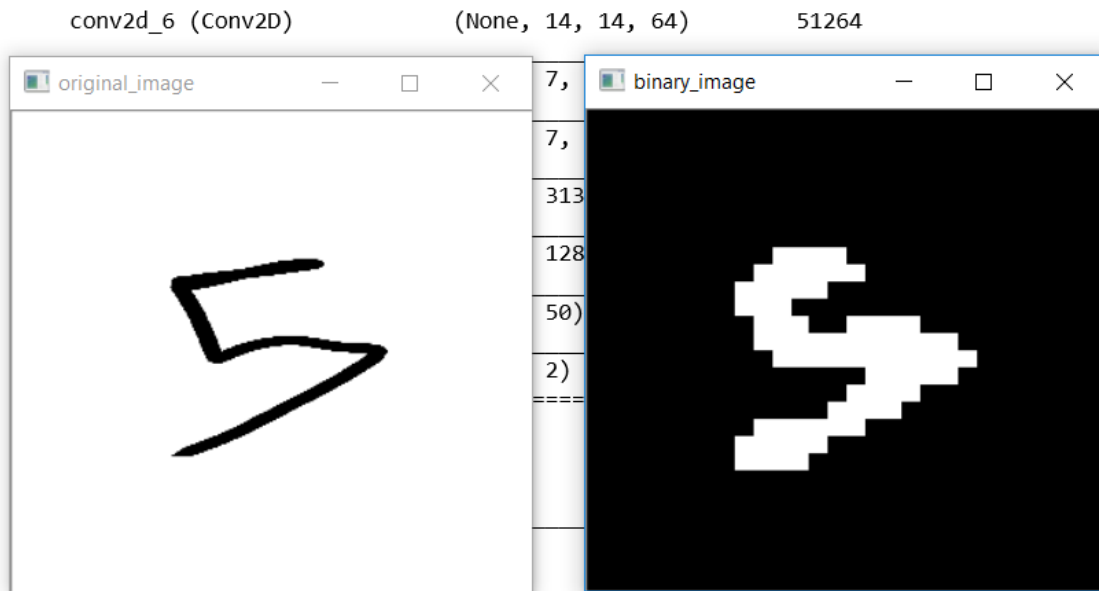
Type the Path of your image file: C:\Users\Herick\Desktop\MS\CS512 Computer Vision\Ass
C:\Users\Herick\Desktop\MS\CS512 Computer Vision\Assignments\Numbers\Sketches\F5.jpeg
True

Type the Path of your image file: C:\Users\Herick\Desktop\MS\CS512 Computer Vision\Ass
C:\Users\Herick\Desktop\MS\CS512 Computer Vision\Assignments\Numbers\Sketches\F5.jpeg
True

Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 28, 28, 32)	832
max_pooling2d_5 (MaxPooling2)	(None, 14, 14, 32)	0
conv2d_6 (Conv2D)	(None, 14, 14, 64)	51264
max_pooling2d_6 (MaxPooling2)	(None, 7, 7, 64)	0
dropout_3 (Dropout)	(None, 7, 7, 64)	0
flatten_3 (Flatten)	(None, 3136)	0
dense_7 (Dense)	(None, 128)	401536
dense_8 (Dense)	(None, 50)	6450
dense_9 (Dense)	(None, 2)	102
Total params: 460,184		
Trainable params: 460,184		
Non-trainable params: 0		

```
[0]  
The digit is odd  
-----
```

However, the below '5' was incorrectly classified as even:



Enter q if you want to exit else enter any other key to continue. h
Type the Path of your image file: C:\Users\Herick\Desktop\MS\CS512 Computer Vision\Assi
C:\Users\Herick\Desktop\MS\CS512 Computer Vision\Assignments\Numbers\Sketches\Five.png
True

Enter q if you want to exit else enter any other key to continue. h
Type the Path of your image file: C:\Users\Herick\Desktop\MS\CS512 Computer Vision\Ass
C:\Users\Herick\Desktop\MS\CS512 Computer Vision\Assignments\Numbers\Sketches\Five.png
True

Layer (type)	Output Shape	Param #
=====		
conv2d_7 (Conv2D)	(None, 28, 28, 32)	832
max_pooling2d_7 (MaxPooling2	(None, 14, 14, 32)	0
conv2d_8 (Conv2D)	(None, 14, 14, 64)	51264
max_pooling2d_8 (MaxPooling2	(None, 7, 7, 64)	0
dropout_4 (Dropout)	(None, 7, 7, 64)	0
flatten_4 (Flatten)	(None, 3136)	0
dense_10 (Dense)	(None, 128)	401536
dense_11 (Dense)	(None, 50)	6450
dense_12 (Dense)	(None, 2)	102
=====		
Total params: 460,184		
Trainable params: 460,184		
Non-trainable params: 0		

[1]
The digit is even

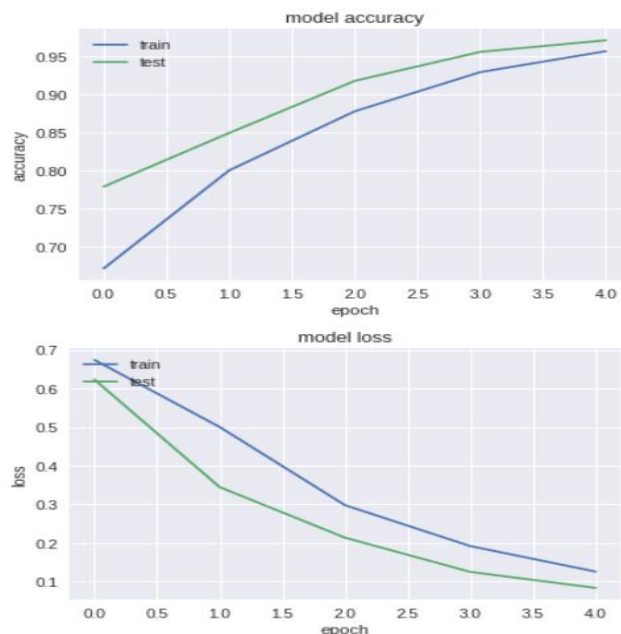
Deliverable 2:

A. Changing the network architecture.

Following architecture was designed.

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 28, 28, 32)	832
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_2 (Conv2D)	(None, 14, 14, 64)	51264
max_pooling2d_2 (MaxPooling2D)	(None, 7, 7, 64)	0
conv2d_3 (Conv2D)	(None, 7, 7, 128)	204928
max_pooling2d_3 (MaxPooling2D)	(None, 3, 3, 128)	0
conv2d_4 (Conv2D)	(None, 3, 3, 256)	819456
max_pooling2d_4 (MaxPooling2D)	(None, 1, 1, 256)	0
dropout_1 (Dropout)	(None, 1, 1, 256)	0
flatten_1 (Flatten)	(None, 256)	0
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 50)	6450
dense_3 (Dense)	(None, 2)	102
Total params: 1,115,928		
Trainable params: 1,115,928		
Non-trainable params: 0		

The above architecture resulted in increase in the number of trainable parameters compared to custom CNN model. Also its validation loss and accuracy improved as shown below.



```
[21] # load the weights that yielded the best validation accuracy
model.load_weights('model.weights.best.hdf5')

# evaluate and print test accuracy
score = model.evaluate(X_test, new_Y_test, verbose=0)
print('Test loss: ', score[0])
print('Test accuracy:', score[1])
```

```
Test loss: 0.08361626062840223
Test accuracy: 0.9705
```

B. Changing Receptive field and strides of custom CNN model.

The Receptive field was set to 3 and stride was set to 2 as shown below.


```
model = Sequential()
model.add(Conv2D(filters = 32, kernel_size = 3, strides = 2, padding = 'same', activation = 'relu', input_shape = ip_shape))
model.add(MaxPooling2D(pool_size = 2, strides = 2))
#model.add(Dropout(0.4))
model.add(Conv2D(filters = 64, kernel_size = 3, strides = 2, padding = 'same', activation = 'relu'))
model.add(MaxPooling2D(pool_size = 2, strides = 2))
model.add(Dropout(0.4))
model.add(Flatten())
model.add(Dense(128, activation = 'relu'))
model.add(Dense(50, activation='relu'))
model.add(Dense(2, activation = 'softmax'))

model.summary()
```

This resulted in decrease in number of parameters compared to custom CNN.

Layer (type)	Output Shape	Param #
conv2d_9 (Conv2D)	(None, 14, 14, 32)	320
max_pooling2d_9 (MaxPooling2	(None, 7, 7, 32)	0
conv2d_10 (Conv2D)	(None, 4, 4, 64)	18496
max_pooling2d_10 (MaxPooling	(None, 2, 2, 64)	0
dropout_3 (Dropout)	(None, 2, 2, 64)	0
flatten_3 (Flatten)	(None, 256)	0
dense_7 (Dense)	(None, 128)	32896
dense_8 (Dense)	(None, 50)	6450
dense_9 (Dense)	(None, 2)	102
Total params: 58,264		
Trainable params: 58,264		
Non-trainable params: 0		

Also, accuracy and loss degraded compared to custom CNN model.




```
[26] # load the weights that yielded the best validation accuracy
      model.load_weights('model.weights.best.hdf5')

      # evaluate and print test accuracy
      score = model.evaluate(X_test, new_Y_test, verbose=0)
      print('Test loss: ', score[0])
      print('Test accuracy:', score[1])
```

```
➡ Test loss: 0.35569756093025207
   Test accuracy: 0.8472
```

C. Changing the optimizer from gradient descent to Adam (lr=0.001)

The accuracy and loss improved compared to custom CNN model.



```
[36] # load the weights that yielded the best validation accuracy
      model.load_weights('model.weights.best.hdf5')

      # evaluate and print test accuracy
      score = model.evaluate(X_test, new_Y_test, verbose=0)
      print('Test loss: ', score[0])
      print('Test accuracy:', score[1])
```


```
➞ Test loss: 0.014513163709636137
   Test accuracy: 0.9951
```

D. Changing various parameters (num filters, dropout = 0.5, lr = 0.01, epoch = 10)
Following architecture is obtained.



Layer (type)	Output Shape	Param #
conv2d_17 (Conv2D)	(None, 28, 28, 16)	416
max_pooling2d_17 (MaxPooling)	(None, 14, 14, 16)	0
conv2d_18 (Conv2D)	(None, 14, 14, 32)	12832
max_pooling2d_18 (MaxPooling)	(None, 7, 7, 32)	0
dropout_7 (Dropout)	(None, 7, 7, 32)	0
flatten_7 (Flatten)	(None, 1568)	0
dense_19 (Dense)	(None, 128)	200832
dense_20 (Dense)	(None, 50)	6450
dense_21 (Dense)	(None, 2)	102
Total params: 220,632		
Trainable params: 220,632		
Non-trainable params: 0		

Number of parameters decreased compared to custom CNN model but accuracy and loss improved.



```
[57] # load the weights that yielded the best validation accuracy
      model.load_weights('model.weights.best.hdf5')

      # evaluate and print test accuracy
      score = model.evaluate(X_test, new_Y_test, verbose=0)
      print('Test loss: ', score[0])
      print('Test accuracy:', score[1])
```

```
Test loss: 0.019207315197959544
Test accuracy: 0.994
```


E. Adding batch normalization

Following results are obtained.



Layer (type)	Output Shape	Param #
=====		
conv2d_7 (Conv2D)	(None, 28, 28, 32)	832
batch_normalization_11 (Batch Normalization)	(None, 28, 28, 32)	128
activation_11 (Activation)	(None, 28, 28, 32)	0
max_pooling2d_7 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_8 (Conv2D)	(None, 14, 14, 64)	51264
batch_normalization_12 (Batch Normalization)	(None, 14, 14, 64)	256
activation_12 (Activation)	(None, 14, 14, 64)	0
max_pooling2d_8 (MaxPooling2D)	(None, 7, 7, 64)	0
dropout_4 (Dropout)	(None, 7, 7, 64)	0
flatten_4 (Flatten)	(None, 3136)	0
dense_10 (Dense)	(None, 128)	401536
batch_normalization_13 (Batch Normalization)	(None, 128)	512
activation_13 (Activation)	(None, 128)	0
dense_11 (Dense)	(None, 50)	6450
batch_normalization_14 (Batch Normalization)	(None, 50)	200
activation_14 (Activation)	(None, 50)	0
dense_12 (Dense)	(None, 2)	102
batch_normalization_15 (Batch Normalization)	(None, 2)	8
activation_15 (Activation)	(None, 2)	0
=====		
Total params: 461,288		
Trainable params: 460,736		
Non-trainable params: 552		

Epoch 00005: val_loss improved from 0.09590 to 0.08429, saving model



```
[38] # load the weights that yielded the best validation accuracy
      model.load_weights('model_weights.best.hdf5')

      # evaluate and print test accuracy
      score = model.evaluate(X_test, new_Y_test, verbose=0)
      print('Test loss: ', score[0])
      print('Test accuracy:', score[1])



➡ Test loss: 0.08429430634379387
   Test accuracy: 0.9841
```

F. Using different weight initializers (He)

Xavier is by default weight initializer in Keras. So, Custom CNN contained Xavier initial weights.

Following result was obtained for He initializer.

1. He_normal:




```
[54] # load the weights that yielded the best validation accuracy
      model.load_weights('model.weights.best.hdf5')

      # evaluate and print test accuracy
      score = model.evaluate(X_test, new_Y_test, verbose=0)
      print('Test loss: ', score[0])
      print('Test accuracy:', score[1])
```

```
☞ Test loss: 0.07527381648942828
   Test accuracy: 0.9724
```

2. He_uniform:

Epoch 00001: val_loss improved from 0.18086 to 0.17011



```
[69] # load the weights that yielded the best validation accuracy
      model.load_weights('model.weights.best.hdf5')

      # evaluate and print test accuracy
      score = model.evaluate(X_test, new_Y_test, verbose=0)
      print('Test loss: ', score[0])
      print('Test accuracy:', score[1]).
```

```
Test loss: 0.06012772022634744
Test accuracy: 0.9784
```

References:

https://docs.opencv.org/3.4.0/d7/d4d/tutorial_py_thresholding.html

<https://stackoverflow.com/questions/42112260/how-do-i-use-the-tensorboard-callback-of-keras>

<https://keras.io/optimizers/>

<https://keras.io/initializers/>

<https://keras.io/callbacks/>

https://github.com/CHIKKIZZY/cs512_TA_CodeSamples/blob/master/AS4/train_with_cola_lab.md